

# Stream Analytics with SQL on Apache Flink®



Fabian Hueske  
@fhueske

**dataArtisans**

Strata Data Conference, London  
*May, 24<sup>th</sup> 2017*

# About me

---



- Apache Flink PMC member
  - Contributing since day 1 at TU Berlin
  - Focusing on Flink's relational APIs since 1.5 years
- Co-founder of data Artisans
- Co-author of "Stream Processing with Apache Flink"
  - Work in progress...
- PhD in Computer Science



# dataArtisans



Original creators of  
**Apache Flink®**



PLATFORM

Providers of the  
**dA Platform**, a supported  
Flink distribution

# Apache Flink



- Platform for scalable stream processing
- Fast
  - Low latency and high throughput
- Accurate
  - Stateful streaming processing in event time
  - Exactly-once state guarantees
- Reliable
  - Highly available cluster setup
  - Snapshot and restart applications



# Powered by Flink



# Flink's DataStream API

---



- The DataStream API is very expressive
  - Application logic implemented as user-defined functions
  - Windows, triggers, evictors, state, timers, async calls, ...
- Many applications follow similar patterns
  - Do not require the expressiveness of the DataStream API
  - Can be specified more concisely and easily with a DSL
- Q: What's the most popular DSL for data processing?

# SQL!



- Many good reasons to use SQL
  - Declarative specification
  - Effective optimization
  - Efficient execution
  - “Everybody” knows SQL
- But does SQL work for streams as well?

# SQL was not designed for streams

---



- Relations are bounded (multi-)sets. ↔ Streams are infinite sequences.
- DBMS can access all data. ↔ Streaming data arrives over time.
- SQL queries return a result and complete. ↔ Streaming queries continuously emit results and never complete.



# Some RDBMS do it anyway 😊

---



- Materialized views (MV) are similar to regular views, but persisted to disk or memory
  - Used to speed-up analytical queries
  - MVs must be updated when the base tables change
- MV maintenance is very similar to SQL on streams
  - Base table updates are a changelog stream
  - MV definition is SQL query to evaluate
  - MV is query result

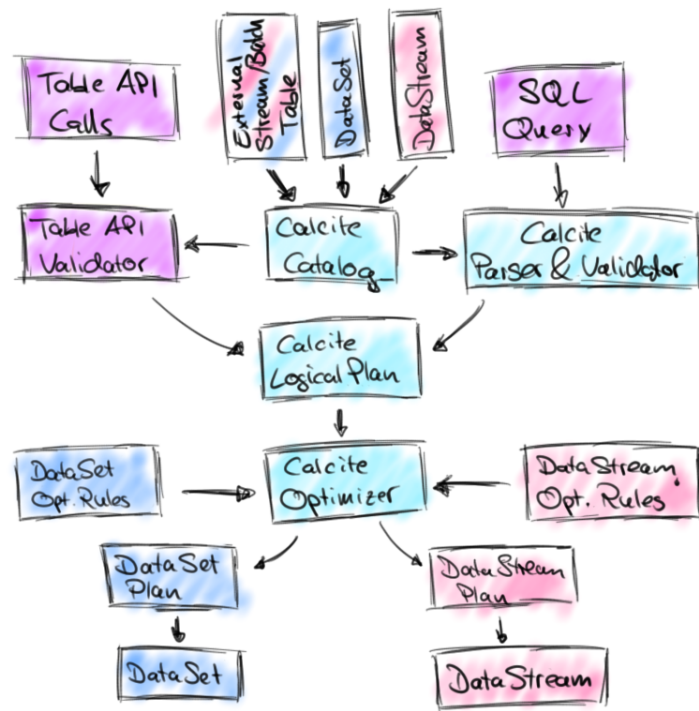
# Apache Flink's Relational APIs



- *Standard SQL & LINQ-style Table API*
- *Unified APIs for batch & streaming data*

***A query specifies exactly the same result regardless whether its input is static batch data or streaming data.***

- Common translation layers
  - Optimization based on Apache Calcite
  - Type system & code-generation
  - Table sources & sinks



# Show me Code!



```
val tableApiResult: Table = tEnv
  .scan("sensors")
  .window(Tumble over 1.hour on 'mtime as 'w)
  .groupBy('w, 'room)
  .select('room, 'w.end, 'temp.avg as 'avgTemp)
```

sensors can be a

- CSV file,
- MySQL table,
- Kafka topic, ...

```
val sqlResult: Table = tEnv.sql("""
| SELECT room,
|       TUMBLE_END(mtime, INTERVAL '1' HOUR),
|       AVG(temp) AS avgTemp
| FROM sensors
| GROUP BY TUMBLE(mtime, INTERVAL '1' HOUR), room
|""").stripMargin)
```

# Continuous Queries



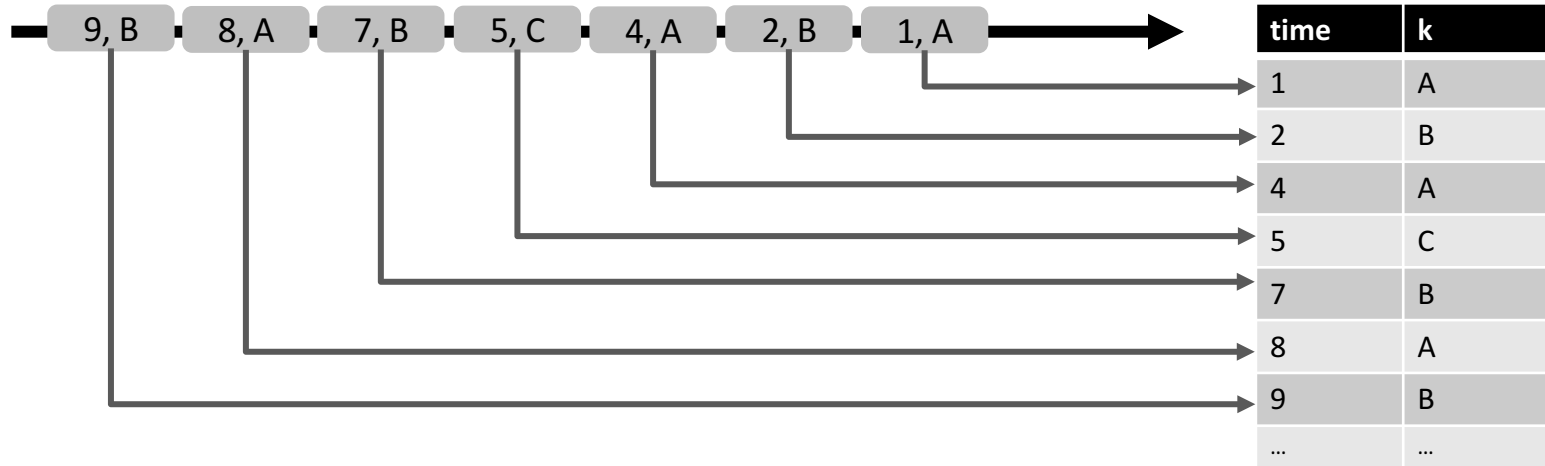
- Core concept is a “*Dynamic Table*”
  - Dynamic tables are changing over time
  - Insert, update, and delete changes
- Dynamic tables can be queries like static batch tables
  - Queries produce new dynamic tables
  - Queries do not terminate
- Stream  $\leftrightarrow$  Dynamic table conversions



# Stream → Dynamic Table



- Stream records are appended to dynamic table
  - Table grows as more data arrives



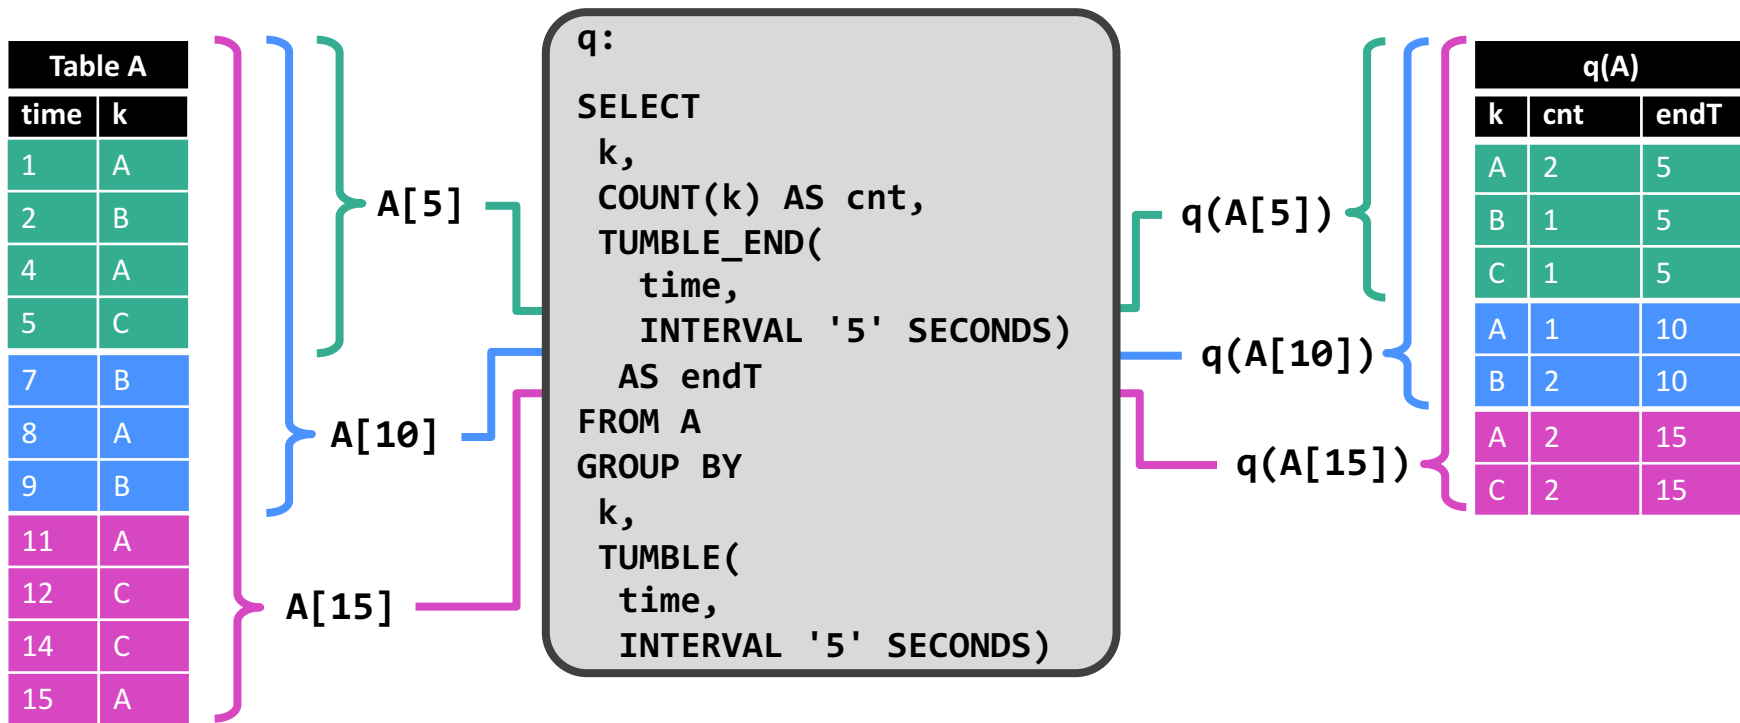
# Querying a Dynamic Table

---

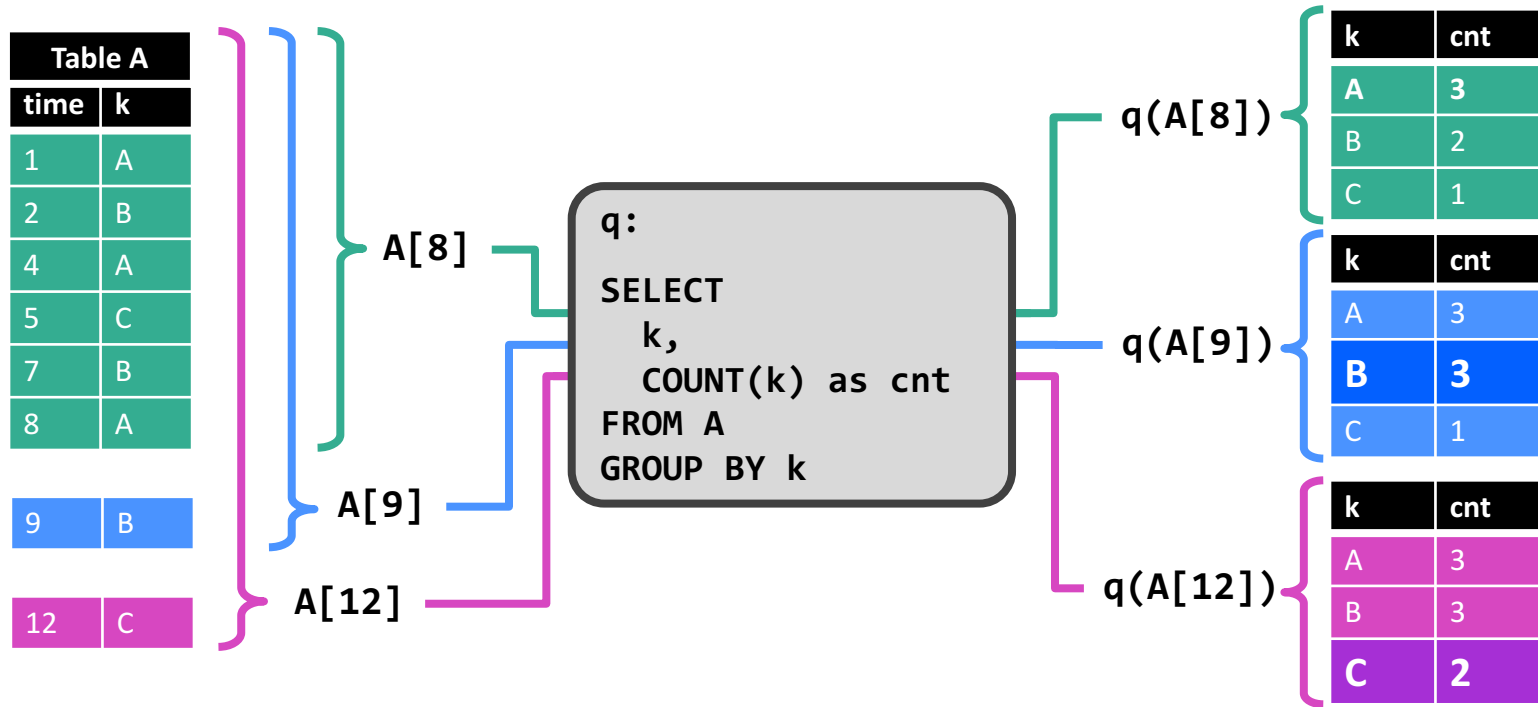


- Dynamic tables change over time
  - $A[t]$ : Table  $A$  at time  $t$
- Dynamic tables are queried with regular SQL
  - $q(A[t])$ : Evaluate query  $q$  on table  $A$  at time  $t$
- As time progresses, the result is continuously updated
  - Similar to maintaining a materialized view

# Querying a Dynamic Table



# Querying a Dynamic Table





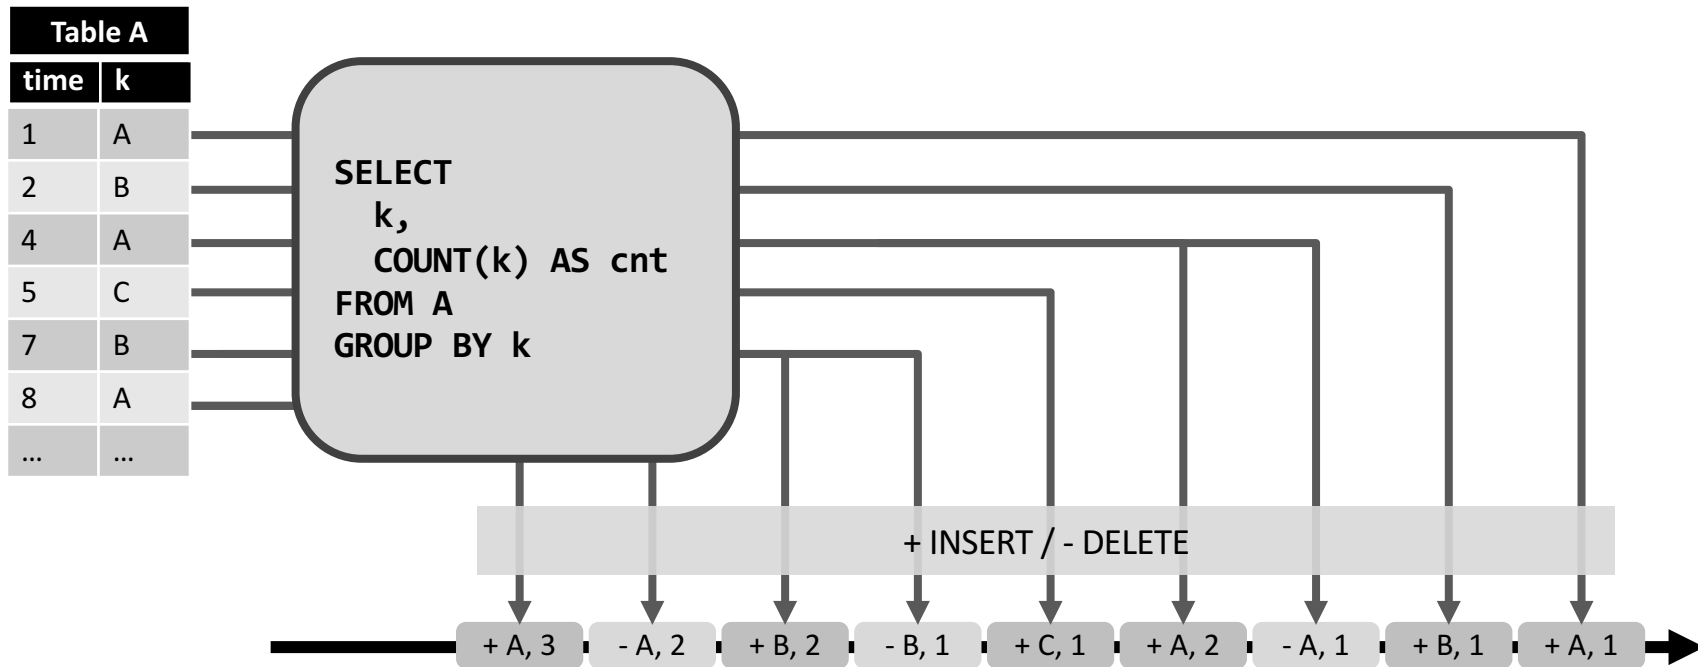
# Dynamic Table → Stream

---

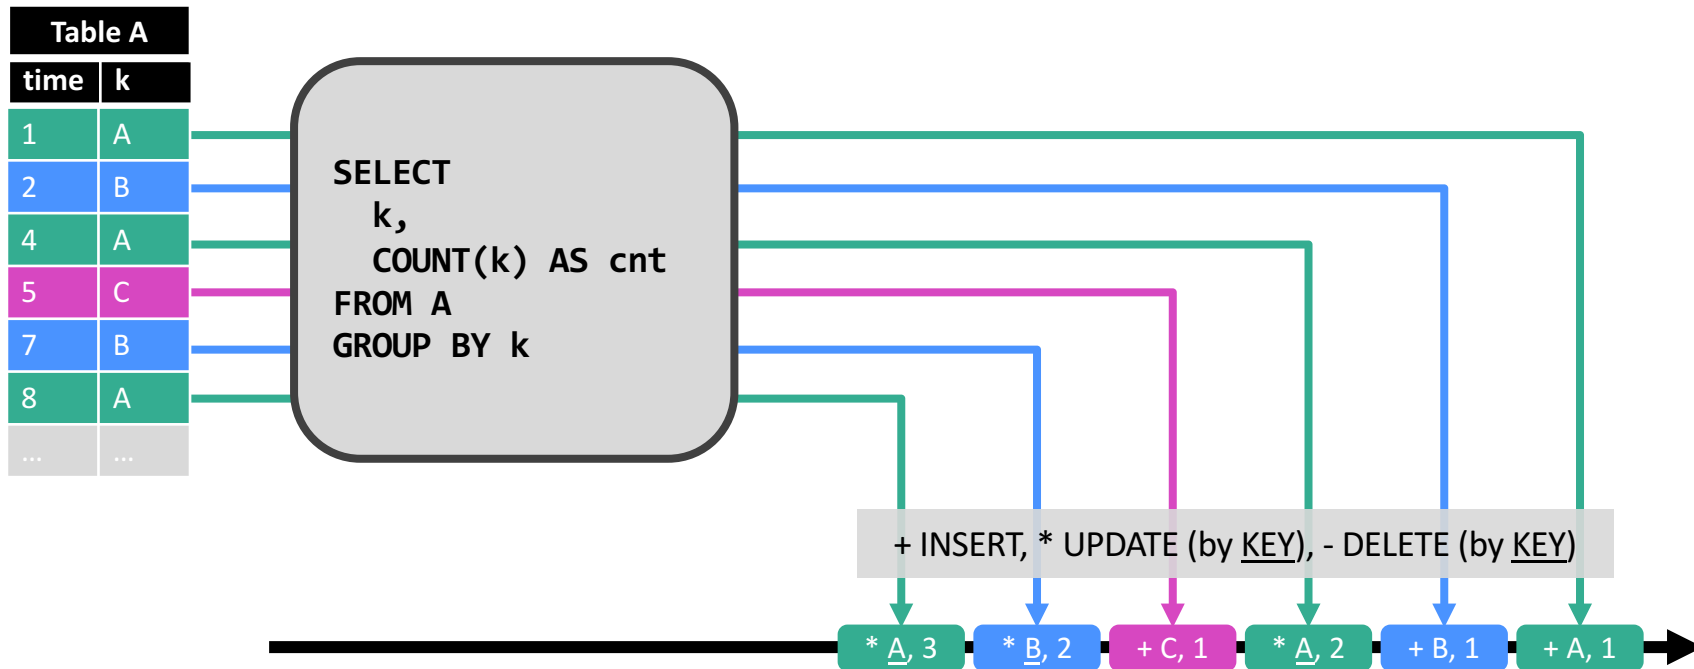


- Converting a dynamic table into a stream
  - Dynamic tables can be updated
  - Updates must be encoded in outgoing stream
  
- Conversion of tables to streams inspired by DBMS logs
  - DBMS use logs to restore databases (and tables)
  - Logs are streams of records that encode updates

# Dynamic Table → Stream: REDO/UNDO



# Dynamic Table → Stream: REDO



# Can We Run Any Query on Dynamic Tables?



- No, there are space and computation constraints ☹️
- State size may not grow infinitely as more data arrives

```
SELECT sessionId, count(*) FROM clicks GROUP BY sessionId;
```

- A change of an input table may only trigger a partial re-computation of the result table

```
SELECT userId, RANK() OVER (ORDER BY highScore) FROM users;
```

# Bounding the Size of Query State



- Adapt the semantics of the query

```
SELECT sessionId, COUNT(*) AS clickCnt
FROM clicks
WHERE last(ctime, INTERVAL '1' DAY)
GROUP BY sessionId
```

- Aggregate data of last 24 hours. Discard older data.
- Trade the accuracy of the result for size of state
  - Remove state for keys that became inactive.

# Current State of SQL & Table API




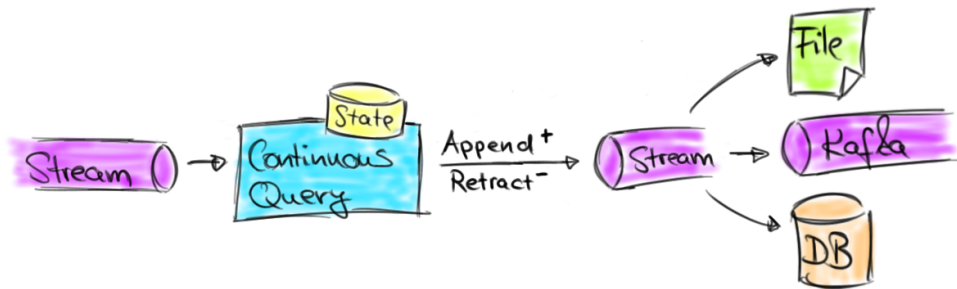
- Flink's relational APIs are rapidly evolving
  - Lots of interest by community and many contributors
  - Used in production at large scale
- Features of the upcoming Flink 1.3.0 release
  - GroupBy & Over windowed aggregates
  - Non-windowed aggregates (with update changes)
  - User-defined aggregation functions



# What can we build with this?



- Continuous ETL & Data Import
- Dashboards & Stateful Microservices



# Conclusion

---



- Table API & SQL support many streaming use cases
  - High-level / declarative specification
  - Automatic optimization and translation
  - Efficient execution
- Updating results enables many exciting applications
  - Materialize a table that is updated by a stream in a KV store
- Check it out!





# FLINK FORWARD

Berlin  
**11-13 Sep 2017**

Flink Forward, the premier conference on Apache Flink®, is coming back to Berlin

Call for Submissions is open

O'REILLY®



# Stream Processing with Apache Flink

FUNDAMENTALS, IMPLEMENTATION, AND OPERATION  
OF STREAMING APPLICATIONS

Fabian Hueske & Vasiliki Kalavri

# Thank you!

@fhueske

@ApacheFlink

@dataArtisans

## Available on O'Reilly Early Release!



We are hiring!

[data-artisans.com/careers](https://data-artisans.com/careers)

